Version info: Code for this page was tested in SAS 9.3.

Multinomial logistic regression is for modeling nominal outcome variables, in which the log odds of the outcomes are modeled as a linear combination of the predictor variables.

Please Note: The purpose of this page is to show how to use various data analysis commands.

It does not cover all aspects of the research process which researchers are expected to do. In

particular, it does not cover data cleaning and checking, verification of assumptions, model

diagnostics and potential follow-up analyses.

Examples of multinomial logistic regression

Example 1. People's occupational choices might be influenced by their parents' occupations and their own education level. We can study the relationship of one's occupation choice with education level and father's occupation. The occupational choices will be the outcome variable which consists of categories of occupations.

Example 2. A biologist may be interested in food choices that alligators make. Adult alligators might have difference preference than young ones. The outcome variable here will be the types of food, and the predictor variables might be the length of the alligators and other environmental variables.

Example 3. Entering high school students make program choices among general program, vocational program and academic program. Their choice might be modeled using their writing score and their social economic status.

Description of the data

For our data analysis example, we will expand the third example using the **hsbdemo** data set. You can download the data here (https://stats.idre.ucla.edu/wp-content/uploads/2016/02/hsbdemo.sas7bdat) .

```
proc contents data = "c:\hsbdemo";
run;
```

The CONTENTS Procedure

Data Set Name	c:\datahsbdemo	Observations	200
Member Type	DATA	Variables	13

```
Engine              V9                                    Indexes                0
Created             Thursday, August 29, 2013 09:42:59 AM Observation Length     40
Last Modified       Thursday, August 29, 2013 09:42:59 AM Deleted Observations   0
Protection                                                Compressed             NO
Data Set Type                                             Sorted                 YES
Label               Written by SAS


Data Representation WINDOWS_64
Encoding            wlatin1  Western (Windows)



          Engine/Host Dependent Information


Data Set Page Size          4096
Number of Data Set Pages    3
First Data Page             1
Max Obs per Page            101
Obs in First Data Page      42
Number of Data Set Repairs  0
Filename                    c:\datahsbdemo.sas7bdat
Release Created             9.0301M1
Host Created                X64_7PRO



      Alphabetic List of Variables and Attributes


 #     Variable    Type    Len     Label

12     AWARDS      Num     3
13     CID         Num     3
 2     FEMALE      Num     3
11     HONORS      Num     3       honores eng
 1     ID          Num     4
 8     MATH        Num     3       math score
 5     PROG        Num     3       type of program
 6     READ        Num     3       reading score
 4     SCHTYP      Num     3       type of school
 9     SCIENCE     Num     3       science score
 3     SES         Num     3
10     SOCST       Num     3       social studies score
 7     WRITE       Num     3       writing score
```

```
 Sort Information


Sortedby        PROG
Validated       YES
Character Set  ANSI
```

The data set contains variables on 200 students. The outcome variable is **prog**, program type. The predictor variables are social economic status, **ses,** a three-level categorical variable and writing score**, write,** a continuous variable. Let's start with getting some descriptive statistics of the variables of interest.

```
proc freq data = "c:\hsbdemo";
tables prog*ses / chisq norow nocol nofreq;
run;
```

```
The FREQ Procedure

Table of PROG by SES

PROG(type of program)     SES

Percent |        1|       2|       3|  Total
--------+--------+--------+--------+
     1 |   8.00 |  10.00 |   4.50 |  22.50
--------+--------+--------+--------+
     2 |   9.50 |  22.00 |  21.00 |  52.50
--------+--------+--------+--------+
     3 |   6.00 |  15.50 |   3.50 |  25.00
--------+--------+--------+--------+
Total        47       95       58      200
           23.50    47.50    29.00   100.00


Statistics for Table of PROG by SES

Statistic                    DF      Value      Prob
---------------------------------------------------------
Chi-Square                    4    16.6044    0.0023
Likelihood Ratio Chi-Square   4    16.7830    0.0021
Mantel-Haenszel Chi-Square    1     0.0598    0.8068
Phi Coefficient                     0.2881
Contingency Coefficient             0.2769
Cramer's V                          0.2037

Sample Size = 200
```

```
proc sort data = "c:\hsbdemo";
by prog;
run;
```

```
proc means data = "c:\hsbdemo";
var write;
by prog;
run;
```

type of program=1

The MEANS Procedure

Analysis Variable : WRITE writing score

N	Mean	Std Dev	Minimum	Maximum
45	51.3333333	9.3977754	31.0000000	67.0000000

type of program=2

Analysis Variable : WRITE writing score

N	Mean	Std Dev	Minimum	Maximum
105	56.2571429	7.9433433	33.0000000	67.0000000

type of program=3

Analysis Variable : WRITE writing score

N	Mean	Std Dev	Minimum	Maximum
50	46.7600000	9.3187544	31.0000000	67.0000000

Analysis methods you might consider

- Multinomial logistic regression: the focus of this page.
- Multinomial probit regression: similar to multinomial logistic regression but with independent normal

error terms.

- Multiple-group discriminant function analysis: A multivariate method for multinomial outcome variables
- Multiple logistic regression analyses, one for each pair of outcomes: One problem with this approach is that each analysis is potentially run on a different sample. The other problem is that without constraining the logistic models, we can end up with the probability of choosing all possible outcome categories greater than 1.
- Collapsing number of categories to two and then doing a logistic regression: This approach suffers from loss of information and changes the original research questions to very different ones.
- Ordinal logistic regression: If the outcome variable is truly ordered and if it also satisfies the assumption of proportional odds, then switching to ordinal logistic regression will make the model more parsimonious.
- Alternative-specific multinomial probit regression: allows different error structures therefore allows to relax the independence of irrelevant alternatives (IIA, see below "Things to Consider") assumption. This requires that the data structure be choice-specific.
- Nested logit model: also relaxes the IIA assumption, also requires the data structure be choice-specific.

Multinomial logistic regression

Below we use **proc logistic** to estimate a multinomial logistic regression model. The outcome **prog** and the predictor **ses** are both categorical variables and should be indicated as such on the **class** statement. We can specify the baseline category for **prog** using (ref = "2") and the reference group for **ses** using (ref = "1"). The **param**=**ref** option on the **class** statement tells SAS to use dummy coding rather than effect coding for the variable **ses**. Note that the levels of **prog** are defined as:

1=general

2=academic (reference group)

3=vocational

```
proc logistic data = "c:\hsbdemo";
class prog (ref = "2") ses (ref = "1") / param = ref;
model prog = ses write / link = glogit;

run;
        The LOGISTIC Procedure

                    Model Information
```

```
Data Set                        c:\datahsbdemo        Written by SAS
Response Variable               PROG                     type of program
Number of Response Levels       3
Model                           generalized logit
Optimization Technique          Newton-Raphson


Number of Observations Read           200
Number of Observations Used           200
```

```
        Response Profile

 Ordered                      Total
   Value          PROG      Frequency

      1             1           45
      2             2          105
      3             3           50
```

Logits modeled use PROG=2 as the reference category.

```
   Class Level Information

                  Design
Class      Value     Variables

SES        1         0      0
           2         1      0
           3         0      1
```

```
            Model Convergence Status

      Convergence criterion (GCONV=1E-8) satisfied.
```

```
      Model Fit Statistics
```

```
                          Intercept
               Intercept       and
Criterion         Only    Covariates

AIC             412.193       375.963
SC              418.790       402.350
-2 Log L        408.193       359.963
```

```
      Testing Global Null Hypothesis: BETA=0


Test                  Chi-Square       DF     Pr > ChiSq
```

Likelihood Ratio	48.2299	6	<.0001
Score	45.1588	6	<.0001
Wald	37.2946	6	<.0001

Type 3 Analysis of Effects

Effect	DF	Wald Chi-Square	Pr > ChiSq
SES	4	10.8162	0.0287
WRITE	2	26.4633	<.0001

Analysis of Maximum Likelihood Estimates

Parameter	PROG		DF	Estimate	Standard Error	Wald Chi-Square	Pr > ChiSq
Intercept	1		1	2.8522	1.1664	5.9790	0.0145
Intercept	3		1	5.2182	1.1635	20.1128	<.0001
SES	2	1	1	-0.5333	0.4437	1.4444	0.2294
SES	2	3	1	0.2914	0.4764	0.3742	0.5407
SES	3	1	1	-1.1628	0.5142	5.1137	0.0237
SES	3	3	1	-0.9827	0.5956	2.7224	0.0989
WRITE	1		1	-0.0579	0.0214	7.3200	0.0068
WRITE	3		1	-0.1136	0.0222	26.1392	<.0001

Odds Ratio Estimates

Effect		PROG	Point Estimate	95% Wald Confidence Limits	
SES	2 vs 1	1	0.587	0.246	1.400
SES	2 vs 1	3	1.338	0.526	3.404
SES	3 vs 1	1	0.313	0.114	0.856
SES	3 vs 1	3	0.374	0.116	1.203
WRITE		1	0.944	0.905	0.984
WRITE		3	0.893	0.855	0.932

- In the output above, the likelihood ratio chi-square of 48.23 with a p-value < 0.0001 tells us that our model as a whole fits significantly better than an empty model (i.e., a model with no predictors)

  - Several model fit measures such as the AIC are listed under Model Fit Statistics

Two models are tested in this multinomial regression, one comparing membership to general versus academic program and one comparing membership to vocational versus academic program. They correspond to the two equations below:

$$ln\left(\frac{P(prog = general)}{P(prog = academic)}\right) = b_{10} + b_{11}(ses = 2) + b_{12}(ses = 3) + b_{13}write$$

$$ln\left(\frac{P(prog = vocation)}{P(prog = academic)}\right) = b_{20} + b_{21}(ses = 2) + b_{22}(ses = 3) + b_{23}write$$

where bs are the regression coefficients.

- A one-unit increase in the variable **write** is associated with a .058 decrease in the relative log odds of being in general program vs. academic program .
- A one-unit increase in the variable **write** is associated with a .1136 decrease in the relative log odds of being in vocation program vs. academic program.
- The relative log odds of being in general program vs. in academic program will decrease by 1.163 if moving from the lowest level of **ses** (**ses**==1) to the highest level of **ses** (**ses**==3).
- The overall effects of **ses** and **write** are listed under "Type 3 Analysis of Effects", and both are significant.

- The ratio of the probability of choosing one outcome category over the probability of choosing the baseline category is often referred to as relative risk (and it is also sometimes referred to as odds as we have just used to described the regression parameters above). Relative risk can be obtained by exponentiating the linear equations above, yielding regression coefficients that are relative risk ratios for a unit change in the predictor variable. In the case of two categories, relative risk ratios are equivalent to odds ratios, which are listed in the output as well.

- The odds ratio for a one-unit increase in the variable **write** is .944 (exp(-.0579) from the regression coefficients above the odds ratios) for being in general program vs. academic program.
- The odds ratio of switching from **ses** = 1 to 3 is .313 for being in general program vs. academic program. In other words, the expected risk of staying in the general program is lower for subjects who are high in **ses**.

Using the **test** statement, we can also test specific hypotheses within or even across logits, such as if the effect of **ses=3** in predicting general versus academic equals the effect of **ses** = 3 in predicting vocational versus academic. Use of the **test** statement requires the unique names SAS assigns each parameter in the model. The option **outest** on the **proc logistic** statement produces an output dataset with the parameter names and values. We can get these names by printing them, and we transpose them to be more readable. The **noobs** option on the **proc print** statement suppresses observation numbers, since they are meaningless in the parameter dataset.

```
proc logistic data = "c:\hsbdemo" outest = mlogit_param;
class prog (ref = "2") ses (ref = "1") / param = ref;
model prog = ses write / link = glogit;
```

```
run;

proc transpose data = mlogit_param;
run;
proc print noobs;
run;
_NAME_           _LABEL_                        PROG

Intercept_1    Intercept: PROG=1               2.852
Intercept_3    Intercept: PROG=3               5.218
SES2_1         SES 2: PROG=1                  -0.533
SES2_3         SES 2: PROG=3                   0.291
SES3_1         SES 3: PROG=1                  -1.163
SES3_3         SES 3: PROG=3                  -0.983
WRITE_1        writing score: PROG=1          -0.058
WRITE_3        writing score: PROG=3          -0.114
_LNLIKE_       Model Log Likelihood          -179.982
```

Here we see the same parameters as in the output above, but with their unique SAS-given names. We are interested in testing whether **SES3_general** is equal to **SES3_vocational**, which we can now do with the test statement. The code preceding the ":" on the **test** statement is a label identifying the test in the output, and it must conform to SAS variable-naming rules (i.e., 32 characters in length or less, letters, numerals, and underscore).

```
proc logistic data = "c:\hsbdemo" outest = mlogit_param;
class prog (ref = "2") ses (ref = "1") / param = ref;
model prog = ses write / link = glogit;
SES3 general vs SES3 vocational: test SES3 1   SES3 3;
```

```
SES3_general_vs_SES3_vocational: test SES3_1 - SES3_3;
run;
```

```
                ***SOME OUTPUT OMITTED***
```

	Linear Hypotheses Testing Results		
Label	Wald Chi-Square	DF	Pr > ChiSq
SES3_general_vs_SES3_vocational	0.0772	1	0.7811

The effect of **ses=3** for predicting general versus academic is not different from the effect of **ses=3** for predicting vocational versus academic.

You can also use predicted probabilities to help you understand the model. You can calculate predicted probabilities using the **lsmeans** statement and the **ilink** option. For multinomial data, **lsmeans** requires **glm** rather than **reference** (dummy) coding, even though they are essentially the same, so be sure to respecify the coding on the **class** statement. However, **glm** coding only allows the last category to be the reference group (**prog** = vocational and **ses** = 3)and will ignore any other reference group specifications.   Below we use **lsmeans** to calculate the predicted probability of choosing program type academic or general at each level of **ses**, holding **write** at its means.

```
proc logistic data = "c:\hsbdemo" outest = mlogit_param;
class prog ses / param = glm;
model prog = ses write / link = glogit;
```

```
lsmeans ses / e llink cl;
run;

***SOME OUTPUT OMITTED***
```

Coefficients for SES Least Squares Means

Parameter	type of program	SES	Row1	Row2	Row3	Row4	Row5	Row6
Intercept	1		1	1	1			
Intercept	2					1	1	1
SES 1	1	1	1					
SES 1	2	1				1		
SES 2	1	2		1				
SES 2	2	2					1	
SES 3	1	3			1			
SES 3	2	3						1
writing score	1		52.775	52.775	52.775			
writing score	2					52.775	52.775	52.775

```
***SOME OUTPUT OMITTED***
```

SES Least Squares Means

type of program	SES	Mean	Standard Error of Mean	Lower Mean	Upper Mean
1	1	0.3582	0.07264	0.2158	0.5006
1	2	0.2283	0.04512	0.1399	0.3168
1	3	0.1785	0.05405	0.07256	0.2844
2	1	0.4397	0.07799	0.2868	0.5925
2	2	0.4777	0.05526	0.3694	0.5861
2	3	0.7009	0.06630	0.5709	0.8309

The predicted probabilities are in the "Mean" column. Thus, for **ses** = 3 and **write** = 52.775, we see that the probability of being the academic program (program type 2) is 0.7009; for the general program (program type 1), the probability is 0.1785.

To obtain predicted probabilities for the program type vocational, we can reverse the ordering of the categories

To obtain predicted probabilities for the program type vocational, we can reverse the ordering of the categories using the **descending** option on the **proc logistic** statement. This will make academic the reference group for **prog** and 3 the reference group for **ses**.

```
proc logistic data = "c:\hsbdemo" outest = mlogit_param descending;
class prog ses / param = glm;
model prog = ses write / link = glogit;
```

```
lsmeans ses / e llink cl;
run;



***SOME OUTPUT OMITTED***
                    Coefficients for SES Least Squares Means
```

Coefficients for SES Least Squares Means

Parameter	type of program	SES	Row1	Row2	Row3	Row4	Row5	Row6
Intercept	3			1	1	1			
Intercept	2						1	1	1
SES 1	3	1	1						
SES 1	2	1					1		
SES 2	3	2			1				
SES 2	2	2						1	
SES 3	3	3				1			
SES 3	2	3							1
writing score	3			52.775	52.775	52.775			
writing score	2						52.775	52.775	52.775

```
***SOME OUTPUT OMITTED***


                    SES Least Squares Means
```

SES Least Squares Means

type of program	SES	Mean	Standard Error of Mean	Lower Mean	Upper Mean
3	1	0.2021	0.05996	0.08459	0.3197
3	2	0.2939	0.05036	0.1952	0.3926
3	3	0.1206	0.04643	0.02960	0.2116
2	1	0.4397	0.07799	0.2868	0.5925
2	2	0.4777	0.05526	0.3694	0.5861
2	3	0.7009	0.06630	0.5709	0.8309

Here we see the probability of being in the vocational program when **ses** = 3 and **write** = 52.775 is 0.1206, which is what we would have expected since (1 – 0.7009 – 0.1785) = 0.1206, where 0.7009 and 0.1785 are the probabilities of being in the academic and general programs under the same conditions.

Things to consider

Things to consider

- The Independence of Irrelevant Alternatives (IIA) assumption: Roughly, the IIA assumption means that adding or deleting alternative outcome categories does not affect the odds among the remaining outcomes.

- Diagnostics and model fit: Unlike logistic regression where there are many statistics for performing model diagnostics, it is not as straightforward to do diagnostics with multinomial logistic regression models. Some model fit statistics are listed in the output.

- Pseudo-R-Squared: The R-squared offered in the output is basically the change in terms of log-likelihood from the intercept-only model to the current model. It does not convey the same information as the R-square for linear regression, even though it is still "the higher, the better".

- Sample size: Multinomial regression uses a maximum likelihood estimation method. Therefore, it requires a large sample size. It also uses multiple equations. Therefore, it requires an even larger sample size than ordinal or binary logistic regression.

- Complete or quasi-complete separation: Complete separation implies that only one value of a predictor variable is associated with only one value of the response variable. You can tell from the output of the regression coefficients that something is wrong. You can then do a two-way tabulation of the outcome variable with the problematic variable to confirm this and then rerun the model without the problematic variable.

- Empty cells or small cells: You should check for empty or small cells by doing a crosstab between categorical predictors and the outcome variable. If a cell has very few cases (a small cell), the model may become unstable or it might not run at all.

- Sometimes observations are clustered into groups (e.g., people within families, students within classrooms). In such cases, you may want to see our page on non-independence within clusters (/stata/dae/non-independence-within-clusters/).

See Also

- SAS Annotated Output: Multinomial Logistic Regression (https://stats.idre.ucla.edu/sas/output/multinomial-logistic-regression/)

References

- Hosmer, D. and Lemeshow, S. (2000) Applied Logistic Regression (Second Edition) (/examples/alr2/). New York: John Wiley & Sons, Inc..

- Agresti, A. (1996) <u>An Introduction to Categorical Data Analysis (/examples/icda/)</u>. New York: John Wiley & Sons, Inc.

Click here to report an error on this page or leave a comment

<u>How to cite this page (https://stats.idre.ucla.edu/other/mult-pkg/faq/general/faq-how-do-i-cite-web-pages-and-programs-from-the-ucla-statistical-consulting-group/)</u>